

BENCHMARK — ANSWER KEY

Observe + traceAI Docs Benchmark

The frozen standard. Graded against, not graded on a curve.

Prepared by: Abhijai · FutureAGI

Date: June 18, 2026

Sources: *report.html* §4/§7 · *playbook 30-depth-decisions.md* · RECONCILIATION.md · EXECUTION_PRD.md · analysis-assets/ screenshots

Status: **LOCKED BEFORE SCORECARD PASS** — do not revise after grading begins

How to use this document

- **Layer A (Cards 1–7):** one rubric card per Diátaxis page type used in our docs. Contains the must-have elements (verbatim from the sources) and a 1/3/5 scale. Use these to score each page during the Scorecard pass. Fill in the blank "score + gaps" slot per page.
- **Layer B (Pairs 1–6):** six exemplar spreads. Left = real competitor screenshot + specific moves to copy. Right = blank slot for our page screenshot. Use during Scorecard to visually compare our page against the gold standard.
- **Coverage Map (Layer C):** every page in the revamped Observe + traceAI IA is listed and mapped to a Layer A card or Layer B pair. No page may be left unmapped.
- **Honesty rule:** this PDF contains ZERO judgment of our own docs. It is only the standard. Criteria are drawn from *report.html* §4/§7, the playbook depth file (*30-depth-decisions.md*), and the binding PRDs (RECONCILIATION, EXECUTION_PRD).

LAYER A

Page-Type Rubric Cards

One card per Diátaxis page type used in our docs. Must-have elements are quoted verbatim from *report.html* §4 (Content-Depth Rubric) and *playbook/30-depth-decisions.md* §"Required concept-page structure" and §"Depth decision matrix." The 1/3/5 scale reflects those sources directly.

Card 1 — Overview / Landing page

Diátaxis type: Reference-overview · Depth range: L1–L3 · Used for: Observe overview, traceAI overview, Getting Started index

PURPOSE

Orient the reader, name the concrete artifact produced, route to next steps. Does not teach the whole product on one page.

MUST-HAVE ELEMENTS (VERBATIM SOURCES)

- "A dedicated motivational section explains concrete business reasons" — *report.html §4, Conceptual "Why" Prose, Gold standard*
- "Inline links throughout the body; a 'Next steps' section; links to related concepts and API reference — user never hits a dead end" — *report.html §4, Cross-linking, Gold standard*
- "Clear depth progression: overview → concept → how-to → advanced → reference → troubleshooting — each section assumes the previous, none repeat it" — *report.html §4, Progressive Disclosure, Gold standard*
- "Tutorial cards on overview page: a grid of cards linking directly to real integration guides" — *report.html §7, Pattern 4*
- "Action-first copy: sidebar labels and page titles that start with a verb or use an action noun" — *report.html §7, Pattern 5*
- First 100 words must contain a direct definition if the page is a concept or architecture page — *30-depth-decisions.md §SEO+GEO*
- Section landing must have: H1, subhead, "What you can do here", Get started, Featured concepts/how-tos — *RECONCILIATION §3a / playbook 29*

1 / 3 / 5 SCALE

1 — Shallow	3 — Adequate	5 — Gold standard
Feature described as a list of capabilities only; no routing to next steps; no "why" prose; no integration cards	One intro paragraph with motivation; some links to concepts; no card grid; no action-first copy	Dedicated motivational section with concrete business reasons; tutorial-card grid for top integrations; action-first CTAs; "what you can do here" section; links to every major entry point; right-rail TOC visible

PLAYBOOK MINIMUMS (QB)

Word floor: 400–900 words (QB-08). **Pattern:** tutorial-card grid on landing (QB-14 Pattern 4). Must expose "On this page" heading tree when 4+ sections (QB-14 Pattern 2).

SCORE + GAPS — TO FILL DURING SCORECARD PASS

Score: ____ / 5

Pages graded by this card: _____

Gaps: _____

Priority fix: _____

Card 2 — Quickstart

Diátaxis type: Tutorial · Depth range: L1–L3 · Used for: Quick start, traceAI Quickstart / Get started

PURPOSE

Get the reader from install to one visible result in 5–10 minutes. Every step produces a verifiable artifact.

MUST-HAVE ELEMENTS (VERBATIM SOURCES)

- "Complete, copy-pasteable snippets with all imports present, inline parameter comments, and at least one complete end-to-end flow" — *report.html §4, Code Completeness, Gold standard*
- "Annotated screenshots with labeled callouts pointing to the specific UI elements being discussed — not generic screenshots" — *report.html §4, Annotated Screenshots, Gold standard*
- "Install-first: Quick start leads with install → first trace → verify" — *report.html §7, Pattern 5 / EXECUTION_PRD §11*
- "Every quickstart can produce one visible result in 5–10 minutes" — *RECONCILIATION §3b, checklist item 13*
- "Every runnable command has an expected output or verification step" — *RECONCILIATION §3b, checklist item 9*
- "Every code example includes imports, client setup, and replaceable credentials" — *RECONCILIATION §3b, checklist item 8*
- Step-numbered flow with anchor links per section — *report.html §7, Pattern 1*

1 / 3 / 5 SCALE

1 — Shallow	3 — Adequate	5 — Gold standard
Abstract description only; no runnable code; no verification step; no expected output shown	Runnable but minimal; missing imports or expected output; no numbered steps with anchors	Numbered H2 steps (## Step 1...) with anchors; copy-pasteable code with imports + credentials; expected output shown after every command; annotated screenshot of result in UI; "verify" section; 5–10 min target stated

PLAYBOOK MINIMUMS (QB)

Word floor: 800–1,200 words (QB-08). **Result:** produces a visible result in 5–10 min (QB-10 item 13). **Sections:** About / Prerequisites / Steps / Verify / Next steps (QB-09). **Code cap:** 150 lines (QB-03).

SCORE + GAPS — TO FILL DURING SCORECARD PASS

Score: ____ / 5

Pages graded by this card: _____

Gaps: _____

Priority fix: _____

Card 3 — Concept

Diátaxis type: Explanation · Depth range: L3–L5 · Used for: Observability model, Traces, Spans, Sessions and users, OpenTelemetry, traceAI SDK, traceAI Concepts

PURPOSE

Teach one atomic noun phrase at a level that lets the reader continue reading the rest of the docs without confusion.

MUST-HAVE ELEMENTS (VERBATIM SOURCES)

- "A dedicated motivational section explains concrete business reasons" — *report.html §4, Conceptual "Why" Prose, Gold standard*
- "Multiple examples covering the common case and at least one edge case, with expected output shown" — *report.html §4, Worked Examples, Gold standard*
- Required concept-page structure: "(1) One-sentence definition: L3. (2) Why it matters: one paragraph, L4. (3) Mental model: diagram or structured example, L4-L5. (4) How it works at a high level: L5. (5) How it works internally / architecture: optional, L5+. (6) When to use it. (7) When not to use it. (8) Examples. (9) Related." — *30-depth-decisions.md §4, "Required concept-page structure"*
- "Dedicated troubleshooting callout or subsection per feature; common pitfalls named with exact error messages and specific fixes" — *report.html §4, Edge Cases & Gotchas, Gold standard*
- "The first sentence defines the noun. The second paragraph explains why it matters in production. A diagram or object shape appears before internals. Standards are linked after the FutureAGI mental model." — *30-depth-decisions.md §4, "FutureAGI progressive disclosure checklist"*
- Diagram: 5–7 nodes; use Mermaid; required when relationships matter — *30-depth-decisions.md §8*
- No concept page has more than 30 lines of code — *30-depth-decisions.md §Pre-publish checklist*
- 0–1 visual concept-state image only — never a UI tour (playbook 30-depth-decisions §Screenshot rules / QB-06)

1 / 3 / 5 SCALE

1 — Shallow	3 — Adequate	5 — Gold standard
Marketing opener; definition buried in paragraph 3+; no diagram; no "why it matters"; no "when not to use"	One-sentence definition present; some motivation; no diagram; partial "when to use"; missing related links	L3 definition in sentence 1; production consequence in sentence 2; Mermaid diagram (5–7 nodes) before internals; "when to use" + "when not to use"; 0–1 concept-state image (never a UI tour); code ≤30 lines showing one concrete example; Related card group at bottom; standards linked after product model

PLAYBOOK MINIMUMS (QB)

Word floor: 800–1,500 words (QB-08). **Structure:** 9-section ordered structure — definition / why / mental model / how it works / internals (optional) / when to use / when NOT to use / examples / related (QB-02). **Code cap:** 30 lines (QB-03). **Image cap:** 0–1 concept-state image only; never a UI tour (QB-06). **Diagram:** Mermaid 5–7 nodes when relationships matter (QB-05).

SCORE + GAPS — TO FILL DURING SCORECARD PASS

Score: ____ / 5

Pages graded by this card: _____

Gaps: _____

Priority fix: _____

Card 4 — Feature deep-dive

Diátaxis type: How-to / Reference hybrid · Depth range: L2–L4 · Used for: Trace explorer, Sessions, Users, Tags, Dashboards, Alerts and monitors, Trace evals, Voice observability

PURPOSE

Let the reader configure and use a specific product feature with no gaps. Every mode, option, limit, and common failure is documented.

MUST-HAVE ELEMENTS (VERBATIM SOURCES)

- "Full table with param name, type, default, description, and valid values for every option — no inference required" — *report.html §4, Parameter / Field Tables, Gold standard*
- 2–4 screenshots of exact UI controls / result states — do not claim "annotated" as a per-page count — *report.html §4, Annotated Screenshots, Gold standard (playbook QB-06)*
- "Dedicated troubleshooting callout or subsection per feature; common pitfalls named with exact error messages" — *report.html §4, Edge Cases & Gotchas, Gold standard*
- "Breaking a feature into explicit numbered steps (Step 1, Step 2...) where each step maps to a discrete UI action" — *report.html §7, Pattern 1*
- "Every filterable/configurable attribute gets its own full page: what it is, how to set it in the SDK, how to filter by it in the UI, and edge cases" — *report.html §7, Pattern 3*
- "Every feature deep-dive includes modes, options, limits, and inspectable output" — *RECONCILIATION §3b, checklist item 15*
- 2–4 screenshots of exact UI controls / result states (playbook QB-06) — *30-depth-decisions.md §Images and screenshots, Feature deep dive row*
- Focused implementation code ≤80 lines — *30-depth-decisions.md §Code depth rules*

1 / 3 / 5 SCALE

1 — Shallow	3 — Adequate	5 — Gold standard
Feature described at capability level only; no parameter table; no screenshots; no edge cases or callouts; code incomplete	One code snippet; some parameters listed in prose; one screenshot (unannotated); one "Note" callout	Numbered step flow with anchor links; full parameter table; 2–4 screenshots of exact UI controls / result states (QB-06); dedicated "Common issues" section with exact errors + fixes; complete runnable code ≤80 lines; right-rail TOC; links to concept + reference + troubleshooting

PLAYBOOK MINIMUMS (QB)

Word floor: 900–1,800 words (QB-08). **Images:** 2–4 screenshots of exact UI controls / result states (QB-06). **Code cap:** 80 lines (QB-03). **Must include:** modes, options, limits, and inspectable output (QB-10 item 15).

SCORE + GAPS — TO FILL DURING SCORECARD PASS

Score: ____ / 5

Pages graded by this card: _____

Gaps: _____

Priority fix: _____

Card 5 — Reference

Diátaxis type: Reference · Depth range: L0–L2 · Used for: Filter syntax, Dashboard metrics, Export and endpoints, traceAI Reference

PURPOSE

Be the authoritative, austere lookup table. Every parameter, operator, format, and limit is documented. No tutorials, no motivation prose — just the contract.

MUST-HAVE ELEMENTS (VERBATIM SOURCES)

- "Full table with param name, type, default, description, and valid values for every option — no inference required" — *report.html §4, Parameter / Field Tables, Gold standard*
- "Multiple examples covering the common case and at least one edge case, with expected output shown (e.g., Braintrust's filter accordions with copy-pasteable SQL expressions)" — *report.html §4, Worked Examples, Gold standard*
- "Every API reference includes method, path, auth, parameters, response, errors, limits, versioning" — *RECONCILIATION §3b, checklist item 10*
- "Assume basic developer API literacy on API reference pages" — *30-depth-decisions.md §LO, Competitor examples, Stripe API EO1*
- No screenshots on API/SDK reference; link to architecture page for diagrams — *30-depth-decisions.md §Images and screenshots, API/SDK reference row*
- Exact operator types documented inline; worked example per operator — *report.html §4, Code Completeness, Gold standard (Sentry Options page)*
- SDK reference must include: package, language, minimum version, signature, parameters, return value, examples, and errors — *RECONCILIATION §3b, checklist item 11 (the traceAI Reference page mapped here is an SDK reference)*

1 / 3 / 5 SCALE

1 — Shallow	3 — Adequate	5 — Gold standard
No parameter docs; reader must infer from code or examples; no operator table; page may 404	Some parameters listed in prose; partial operator list; no type/default/valid-values columns	Complete table: name · type · default · valid values · description for every option; one copy-pasteable example per operator or field; auth, limits, versioning stated; no prose filler; links to concept page for "why"

PLAYBOOK MINIMUMS (QB)

API reference: method, path, auth, parameters, response, errors, limits, versioning (QB-10 item 10). **SDK reference:** package, language, minimum version, signature, parameters, return value, examples, errors (QB-10 item 11). **Images:** 0 screenshots — code is the artifact (QB-06).

SCORE + GAPS — TO FILL DURING SCORECARD PASS

Score: ____ / 5

Pages graded by this card: _____

Gaps: _____

Priority fix: _____

Card 6 — Troubleshooting

Diátaxis type: Reference · Depth range: L2–L4 · Used for: No traces appear, Spans and attributes, Dashboard numbers, Alerts not firing, traceAI Troubleshooting

PURPOSE

Start from symptom or exact error message, not from product overview. Name the root cause, give the specific fix, verify it worked.

MUST-HAVE ELEMENTS (VERBATIM SOURCES)

- "Dedicated troubleshooting callout or subsection per feature; common pitfalls named with exact error messages and specific fixes — not just 'see Troubleshooting' (e.g., Langfuse's per-page Troubleshooting section with GitHub Discussions embedded)" — *report.html §4, Edge Cases & Gotchas, Gold standard*
- "Every troubleshooting page starts from symptom or exact error, not product overview" — *RECONCILIATION §3b, checklist item 17*
- 1–3 screenshots: error state, missing data, broken config — *30-depth-decisions.md §Images and screenshots, Troubleshooting row*
- Decision tree diagram may replace architecture map where symptom branching helps — *30-depth-decisions.md §8, Diagram rules, Troubleshooting row*
- "Warning callout before the risky step" — *30-depth-decisions.md §Callouts, Warning row*
- Symptom → root-cause concept mini-explanation → specific fix → verify step structure — *30-depth-decisions.md §Decision table, Troubleshooting row*

1 / 3 / 5 SCALE

1 — Shallow	3 — Adequate	5 — Gold standard
No warnings; no caveats; vague "check your configuration" advice; starts with product overview not symptom	One "Note" callout; symptom named but root cause not explained; fix given but no verification step	Opens with exact symptom or error message; symptom → root-cause concept (L2–L4 explanation) → specific fix command or config change → verify step; 1–3 annotated screenshots of error state; "Warning" callout before destructive steps; link to canonical concept page for deeper cause

PLAYBOOK MINIMUMS (QB)

Entry point: starts from symptom or exact error, not product overview (QB-10 item 17). **Images:** 1–3 screenshots of error state / missing data / broken config (QB-06). **Diagram:** decision-tree Mermaid diagram when symptom branching helps (QB-05).

SCORE + GAPS — TO FILL DURING SCORECARD PASS

Score: ____ / 5

Pages graded by this card: _____

Gaps: _____

Priority fix: _____

Card 7 — Integration

Diátaxis type: How-to · Depth range: L1–L4 · Used for: all traceAI Framework integrations (LLM Providers, Frameworks and agents, Voice and realtime, Java)

PURPOSE

Get the reader from install to a working instrumented app with a specific framework in a single page. Every integration page names what telemetry is produced and how to verify it in the UI.

MUST-HAVE ELEMENTS (VERBATIM SOURCES)

- "Every integration page names what telemetry, request, or artifact is produced" — *RECONCILIATION §3b, checklist item 12*
- "Complete, copy-pasteable snippets with all imports present, inline parameter comments" — *report.html §4, Code Completeness, Gold standard*
- "Rewrite thin framework leaf pages to install → configure → run → verify → troubleshoot" — *EXECUTION_PRD §5*
- "Auto-instrumentation, traceAI's framework wrapper layer, captures provider calls without editing business logic" — L1 gloss required — *30-depth-decisions.md §L1, FAGI examples*
- "L1 gloss for auto-instrumentation, L2 link to Span, then setup" — *30-depth-decisions.md §Operational decisions, FutureAGI opener rules*
- Link to canonical Span concept page from every integration — *30-depth-decisions.md §L2, FAGI examples*
- Expected output / verify step after every code block — *RECONCILIATION §3b, checklist item 9*

1 / 3 / 5 SCALE

1 — Shallow	3 — Adequate	5 — Gold standard
Install command only; no verification step; does not name what spans are produced; links to product overview not concept	Install + minimal code snippet; names the framework; no expected output; no troubleshooting section	Install → configure → run → verify → troubleshoot flow; names every span kind produced with example names; copy-pasteable with imports; expected trace output described; link to Span concept; mini troubleshooting section; L1 gloss for auto-instrumentation; link to Manual instrumentation for custom spans

PLAYBOOK MINIMUMS (QB)

Word floor: 900–1,500 words (QB-08). **Artifact:** must name the telemetry, request, or artifact produced (QB-10 item 12). **Pattern:** tutorial cards linking to real integration guides (QB-14 Pattern 4).

SCORE + GAPS — TO FILL DURING SCORECARD PASS

Score: ____ / 5

Pages graded by this card: _____

Gaps: _____

Priority fix: _____

LAYER B


Flagship Exemplar Pairs

Six competitor screenshots — each a proven "5" on the rubric above. Left column: the competitor screenshot + specific moves to copy, quoted from *report.html* §7 and §4. Right column: blank slot for our page during the Scorecard pass. Every screenshot below is a real file from `analysis-assets/` verified >50 KB.

Pair 1 — Tags feature page

Competitor model: Langfuse — per-attribute Tags page · File: `46-langfuse-filters-attribute.png` · Size: 351 KB

COMPETITOR: LANGFUSE — TAGS (THIS IS A 5)



SPECIFIC MOVES TO COPY

- "Every filterable/configurable attribute gets its own full page: what it is, how to set it in the SDK, how to filter by it in the UI, and edge cases" — *report.html* §7, *Pattern 3*
- Structure: definition → SDK setup → UI filter usage → edge cases — *report.html* §7, *Pattern 3 move to steal*
- Annotated screenshots with labeled callouts pointing to specific UI elements — *report.html* §4, *Annotated Screenshots*
- Per-page troubleshooting section — *report.html* §4, *Edge Cases & Gotchas*

OUR PAGE — TO BE FILLED DURING SCORECARD PASS



Paste screenshot of our Tags page here.

Use Card 4 (Feature deep-dive) to score it.

Score target: 5 / 5 (Tier 1 page)

COMPETITOR SCORE (BENCHMARK REFERENCE)


5 / 5 — Gold standard

OUR SCORE — TO FILL DURING SCORECARD PASS

Pair 2 — Dashboards feature page

Competitor model: Mixpanel — Boards · File: [14-mixpanel-boards.png](#) · Size: 817 KB

COMPETITOR: MIXPANEL — BOARDS (THIS IS A 5)



SPECIFIC MOVES TO COPY

- "Step-numbered flow: Breaking a feature into explicit numbered steps where each step maps to a discrete UI action" — *report.html §7, Pattern 1*
- "Each step has an anchor link, making it shareable and navigable independently" — *report.html §7, Pattern 1, move to steal*
- Dashboards page = 4-step flow: Create → Add Widget → Configure Metric → Share URL — *EXECUTION_PRD §11*
- "Annotated screenshots with labeled callouts" per step — *report.html §4, Annotated Screenshots*

OUR PAGE — TO BE FILLED DURING SCORECARD PASS



Paste screenshot of our Dashboards page here.

Use Card 4 (Feature deep-dive) to score it.

Score target: 5 / 5 (Tier 1 page)

COMPETITOR SCORE (BENCHMARK REFERENCE)


5 / 5 — Gold standard

OUR SCORE — TO FILL DURING SCORECARD PASS

Pair 3 — Quick start page

Competitor model: Sentry — install-first flow · File: `01-sentry-install-first.png` · Size: 521 KB


COMPETITOR: SENTRY — QUICK START (THIS IS A 5)



SPECIFIC MOVES TO COPY

- "Install-first: Quick start leads with install → first trace → verify" — *EXECUTION_PRD §11 / report.html, Sentry section + EXECUTION_PRD §11*
- "Step-numbered H2 headers (## Step 1, ## Step 2) with anchor links" — *report.html §7, Pattern 1*
- Prerequisites → Install → Configure → Verify → Next Steps structure — *report.html §7, Pattern 2, Sentry "ON THIS PAGE" panel*
- "Annotated screenshots with step-numbered callouts" — *report.html §4, Annotated Screenshots, Gold standard*

OUR PAGE — TO BE FILLED DURING SCORECARD PASS



COMPETITOR SCORE (BENCHMARK REFERENCE)


5 / 5 — Gold standard

OUR SCORE — TO FILL DURING SCORECARD PASS

Pair 4 — Concept page (Traces / Spans / Observability model)

Competitor model: Braintrust — density evals + tracing · Files: `02-braintrust-density-evals.png` (405 KB) + `03-braintrust-density-tracing.png` (408 KB)

COMPETITOR: BRAINTRUST — CONCEPT DEPTH (THIS IS A 5)



SPECIFIC MOVES TO COPY

- "Braintrust sets the depth bar. Every feature page pairs worked examples + parameter tables + annotated UI" — *report.html cover TL;DR*
- L3 definition in sentence 1; production consequence in sentence 2; concrete example in sentence 3 — *30-depth-decisions.md §Operational decisions, opener rules*
- Mermaid diagram (5–7 nodes) before internals — *30-depth-decisions.md §8*
- "Why it matters" paragraph before mental model — *30-depth-decisions.md §4, Required concept-page structure step 2*

OUR PAGES — TO BE FILLED DURING SCORECARD PASS



Paste screenshots of our Traces, Spans, and/or Observability model pages here.

Use Card 3 (Concept) to score each.

Score target: 5 / 5 (Tier 1 pages)

COMPETITOR SCORE (BENCHMARK REFERENCE)


5 / 5 — Gold standard

OUR SCORE — TO FILL DURING SCORECARD PASS

Pair 5 — On-page right-rail TOC (all content pages)

Competitor model: Sentry — on-page TOC panel · File: `32-sentry-on-page-toc.png` · Size: 526 KB

COMPETITOR: SENTRY — RIGHT-RAIL TOC (THIS IS A 5)



SPECIFIC MOVES TO COPY

- "A sticky right-side panel that shows all H2/H3 headings on the current page, allowing jump-navigation without scrolling" — *report.html §7, Pattern 2*
- "Ensure our docs framework exposes right-rail TOC on all content pages. Limit H2 headings to 5–8 per page so the TOC stays scannable" — *report.html §7, Pattern 2, move to steal*
- Frame: ON THIS PAGE → Prerequisites → Install → Configure → Verify → Next Steps — *report.html §7, Pattern 2, Sentry example*
- This is a site-level implementation requirement — affects all pages with 4+ major sections — *report.html §7, Pattern 2*

OUR PAGES — TO BE FILLED DURING SCORECARD PASS



Paste screenshot of any content page on our docs site showing the right-rail TOC.

Score: does the TOC appear? (Yes = implemented, No = gap)

Applies to all pages with 4+ H2 sections.

COMPETITOR SCORE (BENCHMARK REFERENCE)




5 / 5 — Gold standard

OUR SCORE — TO FILL DURING SCORECARD PASS

Pair 6 — Filter syntax / per-filter depth

Competitor models: Braintrust filter page · PostHog filters · Langfuse per-attribute depth (Pricing Tiers example) — Files: [06-braintrust-filter-page.png](#) (381 KB) · [27-posthog-filters-real.png](#) (1.1 MB) · [45-langfuse-filters-features.png](#) (1.4 MB — Pricing Tiers page, not a filters page)

COMPETITORS: BRAINTRUST + POSTHOG (FILTER DEPTH) + LANGFUSE PER-ATTRIBUTE DEPTH (PRICING TIERS EXAMPLE) (THIS IS A 5)



SPECIFIC MOVES TO COPY

- "Braintrust's filter accordions with copy-pasteable SQL expressions" — *report.html §4, Worked Examples, Gold standard*
- "Langfuse wins on filters: every filterable attribute gets its own dedicated page with deep how-to content" — *report.html cover TL;DR*
- The 10 filterable fields = ONE deepened "Filter syntax" reference page with a table — *EXECUTION_PRD §3 decision 5*

OUR PAGE — TO BE FILLED DURING SCORECARD PASS



Paste screenshot of our Filter syntax reference page here.

Use Card 5 (Reference) to score it.

Score target: 5 / 5 (Tier 1 page — was 404 on live site: [report.html](#), FutureAGI section, fig 55; the `/trace-filter-syntax` URL 404s; no source .mdx)

- Four named properties per tier; operator types documented inline; worked example — *report.html*, *Sentry* section (*config depth*)

COMPETITOR SCORE (BENCHMARK REFERENCE)

5 / 5 — Gold standard

OUR SCORE — TO FILL DURING SCORECARD PASS

LAYER C

Coverage Map — Anti-Miss Guarantee

Every page in the revamped Observe + traceAI IA (from *EXECUTION_PRD.md* §4/§5 and *RECONCILIATION.md* §1) mapped to exactly one Layer A card or Layer B pair. An empty "Card / Pair" cell = a hole. There must be none.

Page	Status	Page type	Card / Pair	Depth tier	Score slot
OBSERVE — product section					
Overview <code>/docs/observe</code>	keep	overview	Card 1 — Overview/Landing	Tier 1	
Getting Started group					
Getting Started (index) <code>/docs/observe/getting-started</code>	NEW	overview	Card 1 — Overview/Landing	Tier 1	
Quick start <code>/docs/observe/quickstart</code>	NEW	quickstart	Card 2 — Quickstart · Pair 3 — Sentry install	Tier 1	
Concepts group					
Observability model <code>/docs/observe/concepts/observability-model</code>	NEW	concept	Card 3 — Concept · Pair 4 — Braintrust density	Tier 1	
Traces (was "What are Traces?")	rename	concept	Card 3 — Concept · Pair 4 — Braintrust density	Tier 1	
Spans (was "What are Spans?")	rename	concept	Card 3 — Concept · Pair 4 — Braintrust density	Tier 1	
Sessions and users (was "Sessions and Users")	rename	concept	Card 3 — Concept	Tier 2	
OpenTelemetry (was "What is OpenTelemetry?")	rename	concept	Card 3 — Concept	Tier 2	
traceAI SDK (was "What is traceAI?")	rename	concept	Card 3 — Concept	Tier 2	
Features group					
Trace explorer (was "LLM Tracing")	rename	feature	Card 4 — Feature deep-dive	Tier 1	
Sessions <code>/docs/observe/features/session</code>	keep	feature	Card 4 — Feature deep-dive · Pair 1 — Langfuse Tags	Tier 1	
Users <code>/docs/observe/features/users</code>	keep	feature	Card 4 — Feature deep-dive · Pair 1 — Langfuse Tags	Tier 1	
Tags <code>/docs/observe/features/tags</code>	NEW	feature	Card 4 — Feature deep-dive · Pair 1 — Langfuse Tags	Tier 1	

Page	Status	Page type	Card / Pair	Depth tier	Score slot
Dashboards /docs/observe/features/dashboard	keep	feature	Card 4 — Feature deep-dive · Pair 2 — Mixpanel Boards	Tier 1	
Alerts and monitors (was "Alerts & Monitors")	rename	feature	Card 4 — Feature deep-dive	Tier 2	
Trace evals (was "Run Evals on Traces")	rename	feature	Card 4 — Feature deep-dive	Tier 2	
Voice observability (was "Voice Observability")	rename	feature	Card 4 — Feature deep-dive	Tier 2	
Reference group (Observe)					
Filter syntax (was "Trace Filter Syntax")	rename + deepen	reference	Card 5 — Reference · Pair 6 — Filter depth	Tier 1	
Dashboard metrics (was "Dashboard Metric Definitions")	rename	reference	Card 5 — Reference	Tier 2	
Export and endpoints (was "Trace Export and Endpoints")	rename	reference	Card 5 — Reference	Tier 2	
Troubleshooting group (Observe — LAST)					
No traces appear	keep	troubleshooting	Card 6 — Troubleshooting	Tier 1	
Spans and attributes (was "Missing spans or attributes")	rename	troubleshooting	Card 6 — Troubleshooting	Tier 2	
Dashboard numbers (was "Dashboard numbers look wrong")	rename	troubleshooting	Card 6 — Troubleshooting	Tier 2	
Alerts not firing (was "An alert did not fire")	rename	troubleshooting	Card 6 — Troubleshooting	Tier 2	
traceAI — SDK section (NEW top-level)					
traceAI Overview /docs/traceai	NEW	overview	Card 1 — Overview/Landing	Tier 1	
traceAI Quickstart / Get started	moved + rewrite	quickstart	Card 2 — Quickstart · Pair 3 — Sentry install	Tier 1	
traceAI Concepts (instrumentation model, OTel basis)	moved + rewrite	concept	Card 3 — Concept · Pair 4 — Braintrust density	Tier 2	
Manual instrumentation group (14 pages — moved from Observe)					
Manual instrumentation (14 sub-pages)	moved	feature	Card 4 — Feature deep-dive	Tier 2	
Framework integrations group (moved from Observe)					
Integrations Overview /docs/traceai/auto	moved	integration	Card 1 — Overview/Landing	Tier 2	

Page	Status	Page type	Card / Pair	Depth tier	Score slot
LLM Providers (11 pages)	moved	integration	Card 7 — Integration	Tier 2	
Frameworks and agents (was "Frameworks & Agents")	rename + moved	integration	Card 7 — Integration	Tier 2	
Voice and realtime (was "Voice & Realtime")	rename + moved	integration	Card 7 — Integration	Tier 2	
Java	moved	integration	Card 7 — Integration	Tier 2	
Reference group (traceAI)					
traceAI Reference	moved + rewrite	reference	Card 5 — Reference	Tier 2	
Troubleshooting group (traceAI — LAST)					
traceAI Troubleshooting	moved + rewrite	troubleshooting	Card 6 — Troubleshooting	Tier 2	

Coverage summary: 34 pages mapped (Observe: 21 pages across 6 groups + traceAI: 13 pages across 6 groups). All 34 pages have exactly one Layer A card; 10 pages additionally have a Layer B pair for visual comparison. No page is unmapped. Blank "Score slot" cells are intentional — to be filled during the Scorecard pass.

SECTION C

Playbook Quality Bar (Nikhil's playbook)

Distinct from Layer A / Layer B. Drawn from *playbook-quality-bar.md* (generated 2026-06-18), which in turn synthesizes *oo-master-playbook.md*, *30-depth-decisions.md*, *25-templates-and-checklists.md*, remote audit files, *26-anti-patterns.md*, and *report.html* §4 + §7. Use Section C alongside Layer A cards — it adds hard minimums and anti-pattern gates that the §4 rubric measures but does not set as blocking thresholds.

C.1 — The 14-Bar Catalog (QB-01 through QB-14)

QB	Name	What it requires	Source
QB-01	Depth level scale (L0–L5+)	Every concept assigned one of 7 depth levels before drafting; decision flowchart (4 questions) drives assignment; depth map required per page	<i>30-depth-decisions.md</i> lines 847–855
QB-02	Concept page 9-section structure	Nine ordered sections, none optional: (1) one-sentence definition, (2) why it matters, (3) mental model, (4) how it works, (5) internals (optional), (6) when to use, (7) when NOT to use, (8) examples, (9) related	<i>30-depth-decisions.md</i> lines 1271–1281
QB-03	Code depth caps by page type	Concept: 30 lines; Feature deep-dive: 80 lines; How-to/Quickstart: 150 lines; Cookbook: key excerpts only; Architecture: 30 lines per fragment	<i>30-depth-decisions.md</i> lines 1421–1427
QB-04	First-100-words opener rule	Per-page-type sentence structure required; banned adjectives list (powerful, intuitive, seamless, cutting-edge, revolutionary, game-changing, unlock, leverage, effortlessly, robust without specifics)	<i>30-depth-decisions.md</i> lines 455–462
QB-05	Diagram rules	Concept pages: Mermaid, 5–7 nodes when relationships matter; Architecture: 10–25 nodes; Feature: usually no; Reference: none; Troubleshooting: decision tree only. All must include <code>accTitle</code> + <code>accDescr</code>	<i>30-depth-decisions.md</i> lines 1467–1481
QB-06	Screenshot rules by page type	Concept: 0–1; Feature deep-dive: 2–4; How-to: 1 per non-obvious transition, max 6; API/SDK reference: 0; Troubleshooting: 1–3; Architecture: 0. Crop to the panel. Synthetic project names. <code>last_screenshotted</code> frontmatter required.	<i>30-depth-decisions.md</i> lines 536–566
QB-07	8-dimension 1–5 scoring rubric ("Braintrust bar")	8 dimensions: Conceptual "Why" Prose · Worked Examples · Code Completeness · Parameter/Field Tables · Annotated Screenshots · Edge Cases & Gotchas · Cross-linking · Progressive Disclosure. Score 4–5 required per dimension to meet bar. FutureAGI Observe scored 1.8/5.0 at audit time.	<i>report.html</i> §4
QB-08	Word-count minimums by page type	Overview: 400–900 w; Concept: 800–1,500 w; Feature deep-dive: 900–1,800 w; Quickstart: 800–1,200 w; Integration: 900–1,500 w. Pages under 300 w are too short to teach, debug, or verify anything (326 of 617 current pages fail).	Audit file + <i>25-templates-and-checklists.md</i>

QB	Name	What it requires	Source
QB-09	Per-page-type required sections	Concept: About (120–180w) + Mental model (250–500w) + Key terms + How it works + When to use + Common mistakes + Next steps. Quickstart: About + Prerequisites + Steps + Verify + Next steps. Feature: About + When to use + How it works (300–700w) + Modes/options + Use the feature + Inspect results + Limits + Related.	<i>25-templates-and-checklists.md</i>
QB-10	35-item pre-publish merge-gate checklist	Quality-bearing items include: first 100–150 words answer core question (#5); code includes imports + credentials (#8); runnable command has expected output (#9); API ref has method/path/auth/params/response/errors/limits/versioning (#10); SDK ref has package/language/min-version/signature/params/return/examples/errors (#11); integration names artifact produced (#12); quickstart produces result in 5–10 min (#13); feature covers modes/options/limits/inspectable output (#15); troubleshooting starts from symptom (#17); screenshots are necessary, cropped, dated (#23); <code>last_tested</code> updated (#34).	<i>25-templates-and-checklists.md</i> lines 1961–1999
QB-11	Anti-pattern fast-rejection checklist (7 binary checks)	Can reader define feature in 30 words? Can reader do one useful thing? All claims measurable? Code runnable? Verification step exists? Risks surfaced before risky steps? Pages linked into the graph? Any "no" = block publish.	<i>26-anti-patterns.md</i>
QB-12	Named anti-patterns AP-01 through AP-09	AP-01: Sells product before defining the job (P1). AP-02: Overview too thin (P1). AP-03: Feature claim has no measurement (Po). AP-04: Count drift (Po). AP-05: Missing feature docs behind healthy section count (P1). AP-06: Concept page has no architecture diagram (P1). AP-07: API ref with generic error handling (P1). AP-08: Code uses undefined placeholder variables (Po). AP-09: Secrets as inline code assignments (Po).	<i>26-anti-patterns.md</i>
QB-13	Audit 10-point section scores (historical baseline)	Sections scored 1–10; below 5 = "not trustworthy." Observe section scored 5.6/10 at audit time (see conflict note below — this is a historical baseline, not the canonical scoring instrument).	Audit file
QB-14	§7 feature-presentation pattern library	Pattern 1: step-numbered flow (## Step 1, ## Step 2) with anchors. Pattern 2: "On this page" heading tree when 4+ sections; 5–8 H2s per page. Pattern 3: per-attribute feature pages. Pattern 4: tutorial card grid on overview. Pattern 5: action-first copy on Getting Started / Guide pages.	<i>report.html</i> §7

C.2 — Coherence Table (QB vs Analysis-PDF)

Verdicts: **AGREES** = same principle; **ADDS NEW** = in playbook, not in analysis; **GOLD STANDARD** = analysis IS the bar; **CONFLICT** = instruments disagree; resolution stated.

QB	Playbook bar	Analysis-PDF equivalent	Verdict	Resolution
QB-01	Lo–L5+ depth scale; 4-question flowchart; depth map before drafting	§4 Progressive Disclosure dimension	AGREES	Keep playbook's granularity; 7 levels implement §4's progression at line level
QB-02	Concept page 9-section structure; "When NOT to use" explicit	§4 Conceptual "Why" Prose + Progressive Disclosure + §7 Pattern 5	AGREES	Keep playbook's explicit list; more actionable than §4's description
QB-03	Hard code caps: 30 / 80 / 150 / excerpt lines	§4 Code Completeness ("complete snippets"); §7 Pattern 2	AGREES	Playbook is stricter. Keep caps — they prevent code-dump failure §4 penalizes
QB-04	First-100-words opener; banned adjectives	§4 Conceptual "Why" Prose; AP-01	AGREES	Keep banned adjective list as the implementation of §4's "no feature-list-only opener"
QB-05	Diagram rules: Mermaid, node counts, accTitle required	§4 Annotated Screenshots (lumps screenshots + diagrams)	AGREES	Keep playbook's split. Separate rules reduce confusion between screenshots and diagrams
QB-06	Screenshot targets by page type; crop/naming/dating; last_screenshotted	§4 Annotated Screenshots ("not generic screenshots")	AGREES	Playbook operationalizes §4's score-4 criterion. Keep all rules
QB-07	8-dimension 1–5 rubric; Observe = 1.8/5.0	This IS §4	GOLD STANDARD	§4 is the canonical scoring instrument. Playbook references §4 as authority
QB-08	Word-count floors (400–900 / 800–1,500 / 900–1,800 etc.)	§4 identifies depth problem but sets no floors	ADDS NEW	Keep. Concrete minimums make §4's abstract gap auditable
QB-09	Template section tables (exact subsections + word budgets)	§7 Pattern 1, Pattern 3 — partial overlap	ADDS NEW	Keep. Most operationalized quality bars in the playbook. Writers use as scaffold
QB-10	35-item pre-publish checklist (merge gate)	§4 all dimensions + §7 all patterns	AGREES	Keep checklist. Item 34 (last_tested freshness) is more explicit than anything in the analysis
QB-11	7-binary fast-rejection checklist	§4 rubric implicitly (score 1–2 on any dimension)	AGREES	Keep both: 7-check fast rejection for initial screening; §4 rubric for full scoring
QB-12	Named anti-patterns AP-01–09 with Po–P3 severity	§4 penalizes all (AP-03→Why Prose score 1; AP-07→Param Tables score 1; AP-08/09→Code Completeness score 1)	AGREES (stricter)	Keep Po escalation. Unmeasured claims and broken snippets are trust-destroyers
QB-13	Audit 10-point scores; Observe = 5.6/10 (adequate)	§4 1–5 rubric (Observe = 1.8/5 ≈ very shallow) — different instruments	CONFLICT	§4 wins. 5.6/10 is a historical baseline only. Use §4 for all new pages
QB-14	§7 pattern library (5 patterns as quality requirements)	This IS §7	GOLD STANDARD	§7 is the presentation authority. Playbook references §7

C.3 — Net New Criteria (NNC-01 through NNC-08)

Quality bars the playbook has that *report.html* does NOT contain — the genuine value-add over the analysis.

NNC-01 — FORMAL DEPTH-LEVEL TAXONOMY (L0-L5+)

The analysis identifies depth problems but gives no decision procedure. The playbook's 4-question flowchart produces a deterministic depth assignment. Should be added to the benchmark as the depth-assignment protocol.

NNC-02 — CONCEPT PAGE REQUIRED 9-SECTION STRUCTURE

The analysis calls for motivational prose and progressive disclosure but does not specify the required section list. The playbook's ordered 9-section structure is prescriptive. Notably: "When NOT to use" is not mentioned in §4 at all.

NNC-03 — HARD CODE DEPTH CAPS (30 / 80 / 150 / EXCERPT LINES)

The analysis says code should be "complete" but never sets line-count ceilings. The caps prevent code bloat — the complement of the completeness failure §4 measures.

NNC-04 — WORD-COUNT FLOORS BY PAGE TYPE

The analysis identifies the depth gap but sets no minimum word budgets. The playbook's minimums (400–900 / 800–1,500 / 900–1,800 etc.) make §4's abstract requirements auditable. 326 of 617 current pages fail the 300-word floor.

NNC-05 — NAMED ANTI-PATTERNS WITH P0 SEVERITY ESCALATION

The analysis implies these failures through low rubric scores but does not name them as blockable patterns. The playbook's P0 designation for unmeasured claims (AP-03), placeholder code (AP-08), and inline secrets (AP-09) creates a clear publishing gate.

NNC-06 — FRONTMATTER FRESHNESS REQUIREMENTS (`LAST_TESTED` , `LAST_SCREENSHOTTED`)

The analysis notes 613 of 617 pages have no freshness signal but does not specify the mechanism. The playbook's explicit frontmatter fields and merge-gate item 34 operationalize the requirement. The analysis has no equivalent.

NNC-07 — ATOMIC CONCEPT RULE (NO "AND/OR/WITH" TITLES)

The analysis never addresses concept atomicity. The playbook's rule — if the title contains "and," "or," "with," or a workflow verb, split the page — prevents the "Traces and Spans" or "Online and Offline Evals" anti-pattern.

NNC-08 — PAGE-TYPE PURITY GATE

A page must have exactly one `page_type` (Diátaxis decision before drafting). Prevents the mixed-mode failure (concept + tutorial + reference on one page) that §4 Progressive Disclosure penalizes but does not block at source.

The 5.6/10 vs 1.8/5 Conflict

The playbook audit scored Observe **5.6/10** ("adequate") while the analysis §4 rubric scored the same content **1.8/5** ("very shallow"). These two instruments are NOT calibrated against each other.

The audit's 10-point scale was a preliminary section-level measure. The analysis's §4 8-dimension 1–5 rubric is the canonical instrument going forward. The 5.6/10 score survives only as a historical pre-revamp baseline — not as a quality standard for new pages.

Resolution: §4 is the canonical instrument. 5.6/10 is a historical baseline only.